

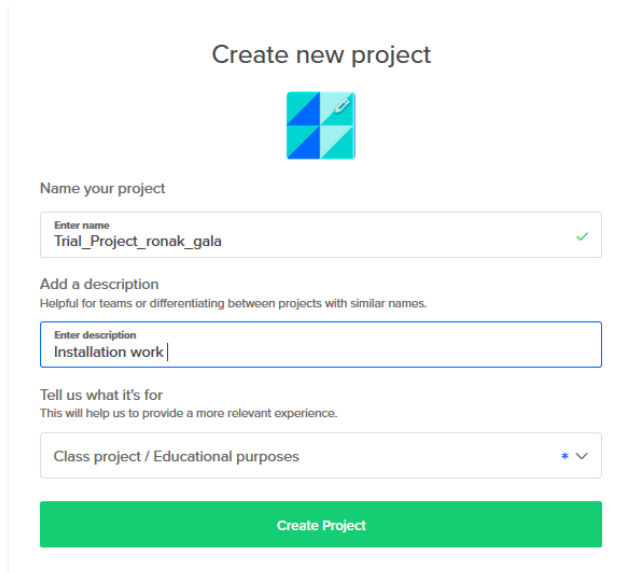
Practical 1 :

Aim : Install, configure and run Hadoop and HDFS and explore HDFS.

Phase 1 : Creating Droplet

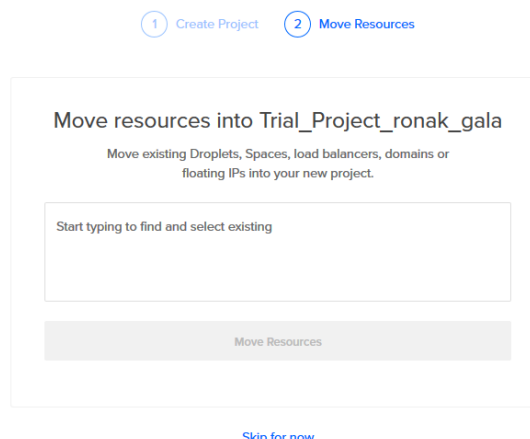
Step 1 : Sign up on <https://www.digitalocean.com/>

Step 2 : Create a new project.



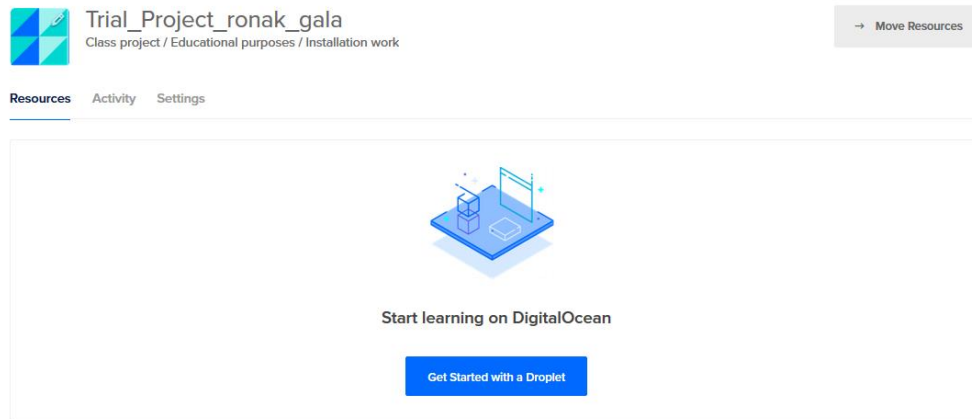
The screenshot shows the 'Create new project' form on DigitalOcean. It includes a title 'Create new project' with a blue icon, a text input for 'Name your project' (filled with 'Trial_Project_ronak_gala'), a text input for 'Add a description' (filled with 'Installation work'), and a dropdown menu for 'Tell us what it's for' (set to 'Class project / Educational purposes'). A green 'Create Project' button is at the bottom.

Step 3 : You can skip the move resources page.



The screenshot shows the 'Move resources into Trial_Project_ronak_gala' page. It has a header with two steps: '1 Create Project' and '2 Move Resources'. The main content area has a text input for 'Start typing to find and select existing' and a greyed-out 'Move Resources' button. A 'Skip for now' link is at the bottom.

Step 4 : Click ‘ Get Started with a Droplet’.



Step 5 : Select the following configurations :-

Choose an Image – Ubuntu

Choose datacentre region – Bangalore

Authentication – OTP (you can also select SSH but you might need to set it up)


How many droplets – 1

Choose a new hostname or go with the provided one


Then select Create Droplet.

Phase 2: Installing JDK

Step 6 : Click on the droplet after it has been created.



Step 7 : Click on console.



A new window will open with the console.

Step 8 : Login in to your registered email id and look for a mail with username and password of your new droplet. Insert the credentials and change the password as per your requirement.

Note : For password the cursor doesn't move, i.e. it won't show the password length just like ubuntu and other Linux flavours.

Step 9 : Check for updates

Command : `sudo apt update`

```
root@ubuntu-s-1vcpu-1gb-blr1-01-trial:~# sudo apt update
Get:1 http://mirrors.digitalocean.com/ubuntu bionic InRelease [242 kB]
Get:2 http://security.ubuntu.com/ubuntu bionic-security InRelease [88.7 kB]
Get:3 http://mirrors.digitalocean.com/ubuntu bionic-updates InRelease [88.7 kB]
Get:4 http://mirrors.digitalocean.com/ubuntu bionic-backports InRelease [74.6 kB]
]
Fetched 494 kB in 2s (271 kB/s)
Reading package lists... Done
Building dependency tree
Reading state information... Done
55 packages can be upgraded. Run 'apt list --upgradable' to see them.
root@ubuntu-s-1vcpu-1gb-blr1-01-trial:~# _
```

Step 10 : Install OpenJDK

Command : `sudo apt install default-jdk`

It will ask for permission as it will take up 727mb of disk, so allow it accordingly.

Step 11 : Check for the java version.

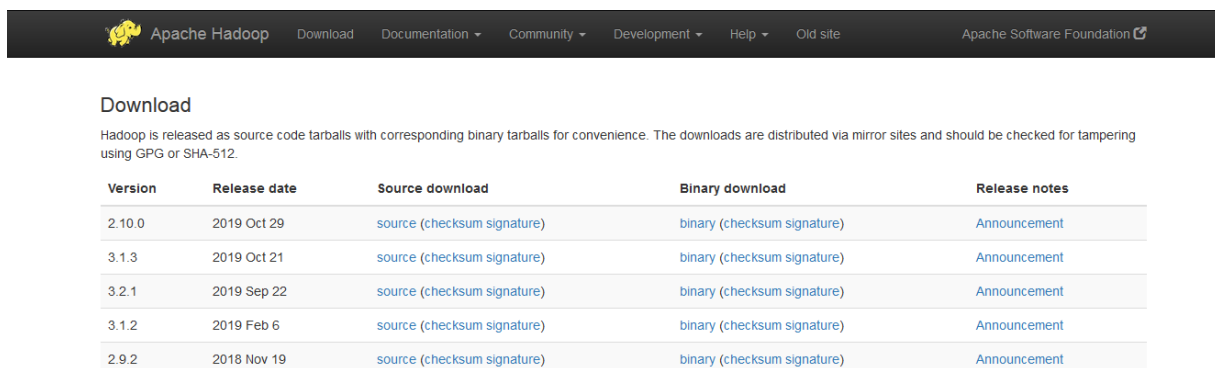
Command : `java -version`

```
root@ubuntu-s-1vcpu-1gb-blr1-01-trial:~# java -version
openjdk version "11.0.6" 2020-01-14
OpenJDK Runtime Environment (build 11.0.6+10-post-Ubuntu-1ubuntu118.04.1)
OpenJDK 64-Bit Server VM (build 11.0.6+10-post-Ubuntu-1ubuntu118.04.1, mixed mode, sharing)
root@ubuntu-s-1vcpu-1gb-blr1-01-trial:~# _
```

Phase 3: Installing Hadoop

Step 12 : Go to <https://hadoop.apache.org/releases.html> to get the stable release.

Step 13 : Navigate to the binary you wish to install. In this guide we will be installing 3.1.3



Version	Release date	Source download	Binary download	Release notes
2.10.0	2019 Oct 29	source (checksum signature)	binary (checksum signature)	Announcement
3.1.3	2019 Oct 21	source (checksum signature)	binary (checksum signature)	Announcement
3.2.1	2019 Sep 22	source (checksum signature)	binary (checksum signature)	Announcement
3.1.2	2019 Feb 6	source (checksum signature)	binary (checksum signature)	Announcement
2.9.2	2018 Nov 19	source (checksum signature)	binary (checksum signature)	Announcement

Step 14 : Click on binary, from your desired version's row.

Step 15 : Copy the http url.



THE APACHE SOFTWARE FOUNDATION 20TH ANNIVERSARY

COMMUNITY-LED DEVELOPMENT "THE APACHE WAY"

Projects People Community License Sponsors

We suggest the following mirror site for your download:

<https://downloads.apache.org/hadoop/common/hadoop-3.1.3/hadoop-3.1.3.tar.gz>

Other mirror sites are suggested below.

It is essential that you verify the integrity of the downloaded file using the PGP signature (.asc file) or a hash (.md5 or .sha* file).

Please only use the backup mirrors to download KEYS, PGP signatures and hashes (SHA* etc) -- or if no other mirrors are working.

Step 16 : On the droplet's console we will use 'wget' to fetch it.

Command : wget <https://downloads.apache.org/hadoop/common/hadoop-3.1.3/hadoop-3.1.3.tar.gz>

Step 17 : We'll use the tar command with the -x flag to extract, -z to uncompress, -v for verbose output, and -f to specify that we're extracting from a file.

Command : tar -xzvf hadoop-3.1.3.tar.gz

Step 18 : We'll move the extracted files into /usr/local

Command : sudo mv hadoop-3.1.3 /usr/local/Hadoop

Phase 4: Configuring Hadoop's Java Home

Step 19 : We'll find java's default path

Command : readlink -f /usr/bin/java | sed "s:bin/java::"

Output : /usr/lib/jvm/java-11-openjdk-amd64/

Step 20 : You can copy this output to set Hadoop's Java home to this specific version, which ensures that if the default Java changes, this value will not. Alternatively, you can use the readlink command dynamically in the file so that Hadoop will automatically use whatever Java version is set as the system default.

To begin, open hadoop-env.sh:

Command : sudo nano /usr/local/hadoop/etc/hadoop/hadoop-env.sh

Step 21 : Move the cursor till the last line and then to set the value to Java_Home write this command:

Option 1 : Set a Static Value

Command : #export JAVA_HOME=\${JAVA_HOME} (this line is a comment)
export JAVA_HOME=/usr/lib/jvm/java-11-openjdk-amd64/

Option 2 :Use Readlink to Set the Value Dynamically

Command : #export JAVA_HOME=\${JAVA_HOME} (this line is a comment)
export JAVA_HOME=\$(readlink -f /usr/bin/java | sed "s:bin/java::")

Press ctrl + X to exit file

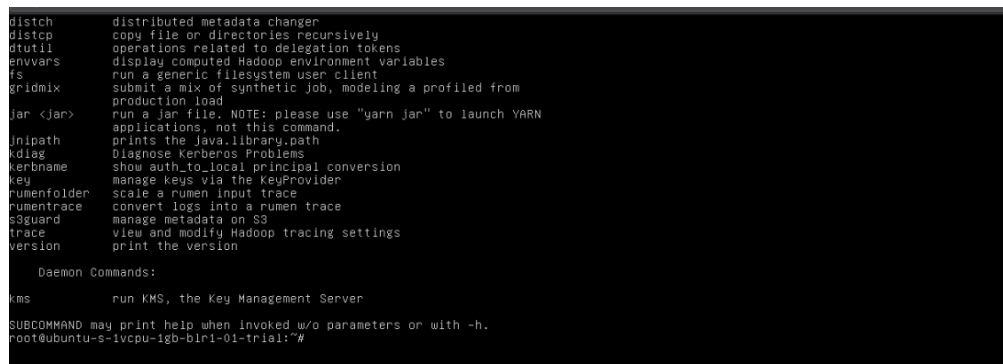
It will ask whether you want to save the modified buffer, answer with YES

It will ask for a file name to write press ENTER

Phase 5: Running Hadoop

Step 22: We will run Hadoop

Command : /usr/local/hadoop/bin/hadoop



```
distch      distributed metadata changer
distcp      copy file or directories recursively
dtutil      operations related to delegation tokens
envvars     display computed Hadoop environment variables
fs          run a generic filesystem user client
gridmix     submit a mix of synthetic job, modeling a profiled from
            production load
jar <jar>    run a jar file. NOTE: please use "yarn jar" to launch YARN
            applications, not this command.
jnipath     prints the java.library.path
kdiag       Diagnose Kerberos Problems
kerbname    show auth_to_local principal conversion
key         manage keys via the KeyProvider
rumenfolder scale a rumen input trace
rumentrace  convert logs into a rumen trace
s3guard     manage metadata on S3
trace       view and modify Hadoop tracing settings
version     print the version

  Daemon Commands:

kms          run KMS, the Key Management Server

SUBCOMMAND may print help when invoked w/o parameters or with -h.
root@ubuntu-s-1vcpu-1gb-blr1-01-trial:~#
```

The help means we've successfully configured Hadoop to run in stand-alone mode.

Step 23 : We'll ensure that it is functioning properly by running the example MapReduce program it ships with. To do so, create a directory called input in our home directory and copy Hadoop's configuration files into it to use those files as our data.

Command 1: `mkdir ~/input`

Command 2: `cp /usr/local/hadoop/etc/hadoop/*.xml ~/input`

```
root@ubuntu-s-1vcpu-1gb-blr1-01-trial:~# mkdir ~/input
root@ubuntu-s-1vcpu-1gb-blr1-01-trial:~# cp /usr/local/hadoop/etc/hadoop/*.xml ~/input
root@ubuntu-s-1vcpu-1gb-blr1-01-trial:~#
```

Step 24 : We can use the following command to run the MapReduce `hadoop-mapreduce-examples` program, a Java archive with several options. We'll invoke its `grep` program, one of the many examples included in `hadoop-mapreduce-examples`, followed by the input directory, `input` and the output directory `grep_example`. The MapReduce `grep` program will count the matches of a literal word or regular expression. Finally, we'll supply the regular expression `allowed[.]*` to find occurrences of the word `allowed` within or at the end of a declarative sentence. The expression is case-sensitive, so we wouldn't find the word if it were capitalized at the beginning of a sentence:

Command : `/usr/local/hadoop/bin/hadoop jar /usr/local/hadoop/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.1.3.jar grep ~/input ~/grep_example 'allowed[.]'`

Note: If the output directory already exists, the program will fail, and rather than seeing the summary, the output will look something like:

Output

...

at java.base/java.lang.reflect.Method.invoke(Method.java:564)

at org.apache.hadoop.util.RunJar.run(RunJar.java:244)

at org.apache.hadoop.util.RunJar.main(RunJar.java:158)

Step 25 : Results are stored in the output directory and can be checked by running cat on the output directory:

Command : cat ~/grep_example/*

```
root@ubuntu-s-1vcpu-1gb-blr1-01-trial:~# cat ~/grep_example/*
21      allowed.
1       allowed
root@ubuntu-s-1vcpu-1gb-blr1-01-trial:~# _
```

The MapReduce task found 21 occurrences of the word allowed followed by a period and one occurrence where it was not. Running the example program has verified that our stand-alone installation is working properly and that non-privileged users on the system can run Hadoop for exploration or debugging.